

# FEATURE SCALING WITH SCIKIT-LEARN

10TH MAY 2017 | IN PYTHON | BY BEN ALEX KEEN

## Feature Scaling with scikit-learn

In this post we explore 3 methods of feature scaling that are implemented in scikit-learn.

- `StandardScaler`
- `MinMaxScaler`
- `RobustScaler`
- `Normalizer`

### Standard Scaler

The `StandardScaler` assumes your data is normally distributed with a mean of 0. After scaling, the distribution is now centred around 0, with a standard deviation of 1.

The mean and standard deviation are calculated for the feature and then used to scale the data.

$$\frac{x_i - \text{mean}(x)}{\text{stdev}(x)}$$

If data is not normally distributed, this is not the best scaler to use.

Let's take a look at it in action:

In [1]:

```
import pandas as pd
import numpy as np
from sklearn import preprocessing
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
matplotlib.style.use('ggplot')
```


In [2]:

```
np.random.seed(1)
df = pd.DataFrame({
    'x1': np.random.normal(0, 2, 10000),
    'x2': np.random.normal(5, 3, 10000),
    'x3': np.random.normal(-5, 5, 10000)
})

scaler = preprocessing.StandardScaler()
scaled_df = scaler.fit_transform(df)
scaled_df = pd.DataFrame(scaled_df, columns=['x1', 'x2', 'x3'])

fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(6, 5))

ax1.set_title('Before Scaling')
sns.kdeplot(df['x1'], ax=ax1)
sns.kdeplot(df['x2'], ax=ax1)
sns.kdeplot(df['x3'], ax=ax1)
ax2.set_title('After Standard Scaler')
sns.kdeplot(scaled_df['x1'], ax=ax2)
sns.kdeplot(scaled_df['x2'], ax=ax2)
sns.kdeplot(scaled_df['x3'], ax=ax2)
plt.show()
```



All features are now on the same scale relative to one another.

## Min-Max Scaler

The `MinMaxScaler` is the probably the most famous scaling algorithm for a single feature:

$$\frac{x_i - \min(x)}{\max(x) - \min(x)}$$

It essentially shrinks the range such that the range is now between 0 and 1.

This scaler works better for cases in which the standard scaler might not work well. If the Gaussian or the standard deviation is very small, the min-max scaler is more appropriate.

However, it is sensitive to outliers, so if there are outliers in the data, the `MinMaxScaler` will be affected. Below.

For now, let's see the `min-max` scaler in action

In [3]:

```
df = pd.DataFrame({
    # positive skew
    'x1': np.random.chisquare(8, 1000),
    # negative skew
    'x2': np.random.beta(8, 2, 1000) * 40,
```


```

# no skew
'x3': np.random.normal(50, 3, 1000)
})

scaler = preprocessing.MinMaxScaler()
scaled_df = scaler.fit_transform(df)
scaled_df = pd.DataFrame(scaled_df, columns=['x1', 'x2',

fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(6, 5))
ax1.set_title('Before Scaling')
sns.kdeplot(df['x1'], ax=ax1)
sns.kdeplot(df['x2'], ax=ax1)
sns.kdeplot(df['x3'], ax=ax1)
ax2.set_title('After Min-Max Scaling')
sns.kdeplot(scaled_df['x1'], ax=ax2)
sns.kdeplot(scaled_df['x2'], ax=ax2)
sns.kdeplot(scaled_df['x3'], ax=ax2)
plt.show()

```



Notice that the skewness of the distribution is maintained but the distributions are now on the same scale so that they overlap.

## Robust Scaler

The `RobustScaler` uses a similar method to the Min-Max scaler but it uses the median and quartiles instead of the min and max, so that it is robust to outliers. Therefore it follows

$$\frac{x_i - Q_1(x)}{Q_3(x) - Q_1(x)}$$

For each feature.

Of course this means it is using the less of the data for scaling so it's not scaling the data.

Let's take a look at this one in action on some data with outliers


In [4]:

```
x = pd.DataFrame({
    # Distribution with lower outliers
    'x1': np.concatenate([np.random.normal(20, 1, 1000),
    # Distribution with higher outliers
    'x2': np.concatenate([np.random.normal(30, 1, 1000),
    })

scaler = preprocessing.RobustScaler()
robust_scaled_df = scaler.fit_transform(x)
robust_scaled_df = pd.DataFrame(robust_scaled_df, columns=x.columns)

scaler = preprocessing.MinMaxScaler()
minmax_scaled_df = scaler.fit_transform(x)
minmax_scaled_df = pd.DataFrame(minmax_scaled_df, columns=x.columns)

fig, (ax1, ax2, ax3) = plt.subplots(ncols=3, figsize=(9, 4))
ax1.set_title('Before Scaling')
sns.kdeplot(x['x1'], ax=ax1)
sns.kdeplot(x['x2'], ax=ax1)
ax2.set_title('After Robust Scaling')
sns.kdeplot(robust_scaled_df['x1'], ax=ax2)
sns.kdeplot(robust_scaled_df['x2'], ax=ax2)
ax3.set_title('After Min-Max Scaling')
sns.kdeplot(minmax_scaled_df['x1'], ax=ax3)
sns.kdeplot(minmax_scaled_df['x2'], ax=ax3)
plt.show()
```



Notice that after Robust scaling, the distributions are brought into the same range and remain outside of bulk of the new distributions.

However, in Min-Max scaling, the two normal distributions are kept in the same 1 range.

## Normalizer

The normalizer scales each value by dividing each value by its magnitude across all features.

Say your features were x, y and z Cartesian co-ordinates your scale factor would be:

$$\frac{x_i}{\sqrt{x_i^2 + y_i^2 + z_i^2}}$$

Each point is now within 1 unit of the origin on this Cartesian co-ordinate system.



In [5]:

```
from mpl_toolkits.mplot3d import Axes3D

df = pd.DataFrame({
    'x1': np.random.randint(-100, 100, 1000).astype(float),
    'y1': np.random.randint(-80, 80, 1000).astype(float),
    'z1': np.random.randint(-150, 150, 1000).astype(float),
})
```

```
scaler = preprocessing.Normalizer()
scaled_df = scaler.fit_transform(df)
scaled_df = pd.DataFrame(scaled_df, columns=df.columns)

fig = plt.figure(figsize=(9, 5))
ax1 = fig.add_subplot(121, projection='3d')
ax2 = fig.add_subplot(122, projection='3d')
ax1.scatter(df['x1'], df['y1'], df['z1'])
ax2.scatter(scaled_df['x1'], scaled_df['y1'], scaled_df['z1'])
plt.show()
```



Note that the points are all brought within a sphere that is at most 1 unit from the origin. The axes that were previously different scales are now all one scale.

